

tutorial_2

October 4, 2019

1 Tutorial 2 - Relaxing OLS assumptions

2 Homoskedasticity

2.1 Simulate data

We have a sample $(x_t, y_t)_{t=1}^T$ generated from the following model

$$y_t = \alpha + \beta x_t + e_t$$

where $x \perp e$ and x, e follow, respectively, a $N(0, 1)$ and $N(0, \sigma^2)$ distribution.

```
[1]: import numpy as np                # linear algebra
import statsmodels.formula.api as smf # linear regression
import matplotlib.pyplot as plt       # plotting
import pandas as pd                   # dataframes
from scipy import stats                # N(0, 1) pdf

# simulation parameters
npop = 1000 # must be even
alpha = .5
beta = 1.2
sigma = .8


# generate random data
# np.random.seed(1) # important! Seed for reproducibility
x1 = np.random.normal(0, 1, npop)
e1 = np.random.normal(0, sigma, npop)
y1 = alpha + beta*x1 + e1

# create dataframe with population data
df1 = pd.DataFrame()
df1['y'] = y1
df1['x'] = x1

[2]: # plot scatter
fig, ax = plt.subplots()
ax.set_title("Scatter of $$ on $$")
ax.scatter(df1['x'], df1['y'], s=4)
```

```
ax.set_xlabel('$x$')
ax.set_ylabel('$y$')
```

[2]: Text(0, 0.5, '\$y\$')



2.1.1 True vs empirical distribution OLS estimator

Here is the experiment that we are going to run in a somewhat heuristic way:

1. we fit a linear regression $y \sim x$ on the *whole* sample and save the estimated $\hat{\beta}$ and its SE
2. we sample a random subset of our data $(x_t^{(i)}, y_t^{(i)})_t$
3. we fit the same linear regression $y \sim x$ on this sub-sample and save the estimated $\hat{\beta}^{(i)}$
4. we repeat steps 2-3) for $i = 1, 2, \dots, Nsim$ and obtain an empirical distribution of all the $(\hat{\beta}^{(i)})_{i=1}^N$

Then, if we correctly specified our model, we should expect that the empirical distribution of the $\hat{\beta}^{(i)}$ approaches the “true” asymptotic distribution of $\hat{\beta}$ which - under homoskedasticity - we know is

$$\sqrt{T} \begin{pmatrix} \hat{\alpha} - \alpha \\ \hat{\beta} - \beta \end{pmatrix} \xrightarrow{T \rightarrow +\infty} \mathcal{N} \left(\begin{pmatrix} 0 \\ 0 \end{pmatrix}, \sigma^2 plim \left(\frac{X'X}{T} \right)^{-1} \right)$$

```
[3]: # 1) fit beta hat from full dataset
reg_mod1 = smf.ols("y ~ x", df1)
fit_hom1 = reg_mod1.fit() # homosk errors
betahat_se_hom1 = fit_hom1.bse[1]
fit_hom1.summary()
```

```
[3]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                OLS Regression Results
=====
Dep. Variable:                  y      R-squared:                0.679
Model:                            OLS      Adj. R-squared:           0.678
Method:                 Least Squares      F-statistic:                2109.
Date:                Fri, 04 Oct 2019      Prob (F-statistic):          2.45e-248
Time:                  10:00:41      Log-Likelihood:            -1211.3
No. Observations:                1000      AIC:                       2427.
Df Residuals:                     998      BIC:                       2436.
Df Model:                           1
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.4482	0.026	17.413	0.000	0.398	0.499
x	1.1730	0.026	45.922	0.000	1.123	1.223

```

=====
Omnibus:                        3.957      Durbin-Watson:           2.028
Prob(Omnibus):                  0.138      Jarque-Bera (JB):        4.180
Skew:                           0.083      Prob(JB):                0.124
Kurtosis:                      3.270      Cond. No.:               1.04
=====

```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
```

```
"""
```

```
[4]: # 2-3) fit model on random sub-sample
def bootstrap_lm_hom(x, params):
    """ Fit univariate lin reg to random sample

    The function assumes we know the true DGP

    Arguments
    -----
    x : np.array, univariate explanatory variable
    params : 1D np.array, alpha, beta and sigma

    Returns
```

```

-----
    beta_boot : estimated beta coeff on bootstrap sample
    """
    # generate bootstrap sample
    alpha, beta, sigma = params
    n = len(x)
    indices = np.random.choice(np.arange(n), n)

    # generate residuals using right DGP
    x_boot = x[indices]
    e_boot = np.random.normal(0, sigma, n)
    y_boot = alpha + beta*x_boot + e_boot

    # create dataframe with bootstrapped data
    df = pd.DataFrame()
    df['y'] = y_boot
    df['x'] = x_boot

    # fit linear regression
    fit = smf.ols("y ~ x", df).fit()
    beta_boot = fit.params[1]

    return beta_boot

```

```

[5]: # fit on random sample
nsim = 1000
f1 = lambda i: bootstrap_lm_hom(df1['x'], [alpha, beta, sigma])
betas1 = list(map(f1, range(nsim)))

```

```

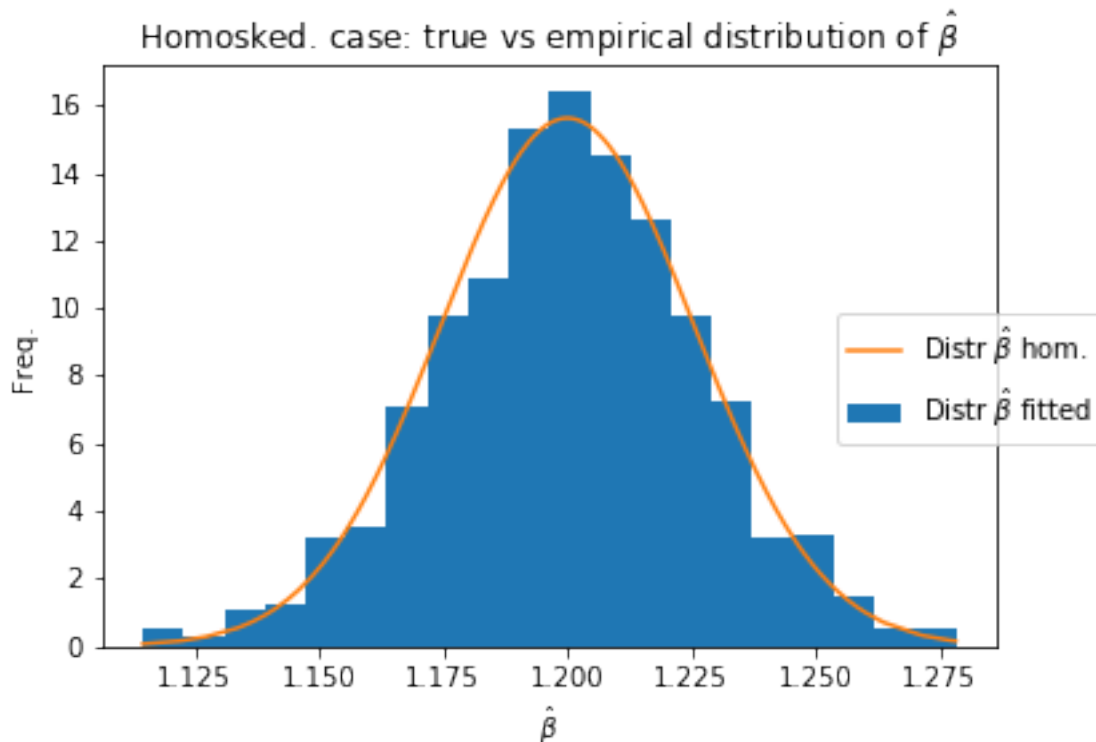
[6]: # 4) plot empirical vs true density of betahat
fig, ax = plt.subplots()
ax.hist(betas1, bins=20, density=True, label="Distr  $\hat{\beta}$  fitted" )
ax.plot(sorted(betas1), stats.norm.pdf(sorted(betas1), beta, betahat_se_hom1),
        label="Distr  $\hat{\beta}$  hom.")
ax.set_xlabel(" $\hat{\beta}$ ")
ax.set_ylabel("Freq.")
ax.set_title("Homosked. case: true vs empirical distribution of  $\hat{\beta}$ ")
fig.legend(loc="center right")

```

```

[6]: <matplotlib.legend.Legend at 0x2f3632e4f28>

```



2.2 Heteroskedasticity

2.2.1 Simulate data

We have a sample $(x_t, y_t)_{t=1}^T$ generated from the following model

$$y_t = \alpha + \beta x_t + e_t$$

where $x \perp e$, x follows a $N(0, 1)$ and $e_t \sim N(0, |1 - x_t|)$ for $t = 1, 2, \dots, T$.

```
[7]: # simulation parameters
npop = 3000
alpha = .5
beta = 1.2


# generate random data
np.random.seed(1234) # important! Seed for reproducibility
x2 = np.random.normal(0, 1, npop)
e2 = np.random.normal(0, np.abs(1 - x2), npop)
y2 = alpha + beta*x2 + e2

# create dataframe with population data
df2 = pd.DataFrame()
df2['y'] = y2
```

```
df2['x'] = x2
```

```
[8]: fig, ax = plt.subplots()
      ax.set_title("Scatter of $y$ on $x$")
      ax.scatter(df2['x'], e2, s=4)
      ax.set_xlabel('$x$')
      ax.set_ylabel('$y$')
```

```
[8]: Text(0, 0.5, '$y$')
```



2.2.2 True vs empirical distribution OLS estimator

Here is the experiment that we are going to run in a somewhat heuristic way:

1. we fit a linear regression $y \sim x$ on the whole sample and save the estimated $\hat{\beta}$ and its SE. We also (wrongly) assume *homoskedastic* errors.
2. we fit a linear regression $y \sim x$ on the whole sample and save the estimated $\hat{\beta}$ and its SE. We also (correctly) assume *heteroskedastic* errors.
3. we sample a random subset of our data $(x_t^{(i)}, y_t^{(i)})_t$
4. we fit the same linear regression $y \sim x$ on this sub-sample and save the estimated $\hat{\beta}^{(i)}$
5. we repeat steps 3-4) for $i = 1, 2, \dots, Nsim$ and obtain an empirical distribution of all the $(\hat{\beta}^{(i)})_{i=1}^N$

Then, if we correctly specified our model, we should expect that the empirical distribution of the $\hat{\beta}^{(i)}$ approaches the “true” asymptotic distribution of $\hat{\beta}$ which - under heteroskedasticity - we know is

$$\sqrt{T} \begin{pmatrix} \hat{\alpha} - \alpha \\ \hat{\beta} - \beta \end{pmatrix} \rightarrow \mathcal{N} \left(\begin{pmatrix} 0 \\ 0 \end{pmatrix}, (X'X/T)^{-1} \frac{X'\Omega X}{T} (X'X/T)^{-1} \right)$$

with

$$\Omega = \begin{pmatrix} \sigma_1^2 & 0 & \dots & 0 \\ 0 & \sigma_2^2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \sigma_T^2 \end{pmatrix}.$$

Note: we have suppressed the plim operator for the sake of simplicity.

```
[9]: # 1) fit beta hat from full dataset under homeskd
reg_mod2 = smf.ols("y ~ x", df2)
fit_hom2 = reg_mod2.fit() # homosk errors
betahat_se_hom2 = fit_hom2.bse[1]

# plot fit vs residuals
fig, ax = plt.subplots()
ax.set_title("Fitted - Residual plot")
ax.scatter(fit_hom2.fittedvalues, fit_hom2.resid, s=4)
ax.set_xlabel('Fitted')
ax.set_ylabel('Residuals')
# Show summary of the regression results
fit_hom2.summary()
```

```
[9]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                OLS Regression Results
=====
Dep. Variable:                  y      R-squared:                0.425
Model:                            OLS      Adj. R-squared:           0.424
Method:                 Least Squares      F-statistic:             2212.
Date:                Fri, 04 Oct 2019      Prob (F-statistic):       0.00
Time:                  10:00:56      Log-Likelihood:          -5218.4
No. Observations:                3000      AIC:                    1.044e+04
Df Residuals:                    2998      BIC:                    1.045e+04
Df Model:                            1
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.5279	0.025	20.964	0.000	0.478	0.577
x	1.2012	0.026	47.035	0.000	1.151	1.251

```
=====
```

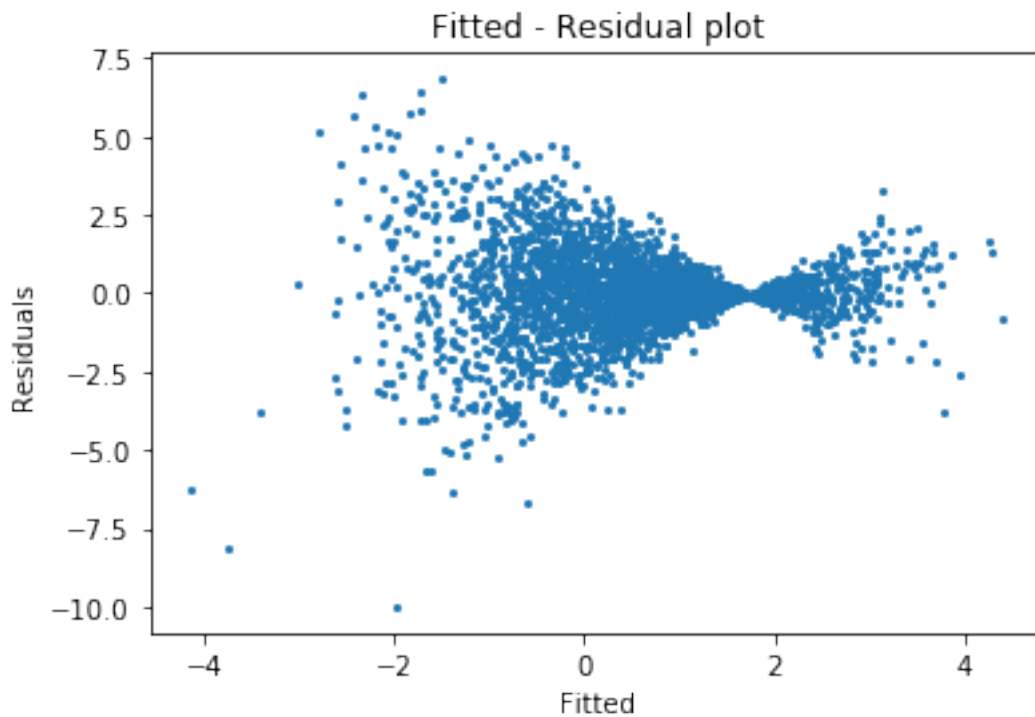
Omnibus:	283.977	Durbin-Watson:	1.971
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1933.630
Skew:	-0.126	Prob(JB):	0.00
Kurtosis:	6.925	Cond. No.	1.04

=====

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""



```
[10]: # 2-3) fit model on random sub-sample
def bootstrap_lm_het(x, params):
    """ Fit univariate lin reg to random sample

    The function assumes we know the true DGP

    Arguments
    -----
    x : np.array, univariate explanatory variable
    params : 1D np.array, alpha, beta and sigma

    Returns
    -----
```



```

        beta_boot : estimated beta coeff on bootstrap sample
        """
        # generate bootstrap sample
        alpha, beta = params
        n = len(x)
        indices = np.random.choice(np.arange(n), n)

        # generate residuals using right DGP
        x_boot = x[indices]
        e_boot = np.random.normal(0, np.abs(1 - x_boot), n)
        y_boot = alpha + beta*x_boot + e_boot

        # create dataframe with bootstrapped data
        df = pd.DataFrame()
        df['y'] = y_boot
        df['x'] = x_boot

        # fit linear regression
        fit = smf.ols("y ~ x", df).fit()
        beta_boot = fit.params[1]

    return beta_boot

```

```

[11]: # fit on random sample
nsim = 1000
np.random.seed(987)
f2 = lambda i: bootstrap_lm_het(df2['x'], [alpha, beta])
betas2 = list(map(f2, range(nsim)))

```

```

[12]: # 2) fit beta hat from full dataset under heterosk
fit_het2 = reg_mod2.fit(cov_type = 'HC0') # White heter errors
betahat_se_het2 = fit_het2.bse[1]
print(f"beta SE hom {betahat_se_hom2}")
print(f"beta SE het {betahat_se_het2}")
fit_het2.summary()

```

beta SE hom 0.025538200266201246

beta SE het 0.036871988227998316

```

[12]: <class 'statsmodels.iolib.summary.Summary'>
"""

```

```

                                OLS Regression Results
=====
Dep. Variable:                  y      R-squared:                0.425
Model:                            OLS      Adj. R-squared:           0.424
Method:                 Least Squares      F-statistic:             1061.
Date:                 Fri, 04 Oct 2019      Prob (F-statistic):       1.44e-199
Time:                 10:01:09      Log-Likelihood:          -5218.4

```

```

No. Observations:      3000    AIC:      1.044e+04
Df Residuals:          2998    BIC:      1.045e+04
Df Model:              1
Covariance Type:      HCO

```

```

=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
Intercept      0.5279      0.026     20.260      0.000      0.477      0.579
x              1.2012      0.037     32.577      0.000      1.129      1.273
=====
Omnibus:                283.977    Durbin-Watson:                1.971
Prob(Omnibus):           0.000    Jarque-Bera (JB):            1933.630
Skew:                   -0.126    Prob(JB):                     0.00
Kurtosis:                6.925    Cond. No.                     1.04
=====

```

Warnings:

```

[1] Standard Errors are heteroscedasticity robust (HCO)
"""

```

```

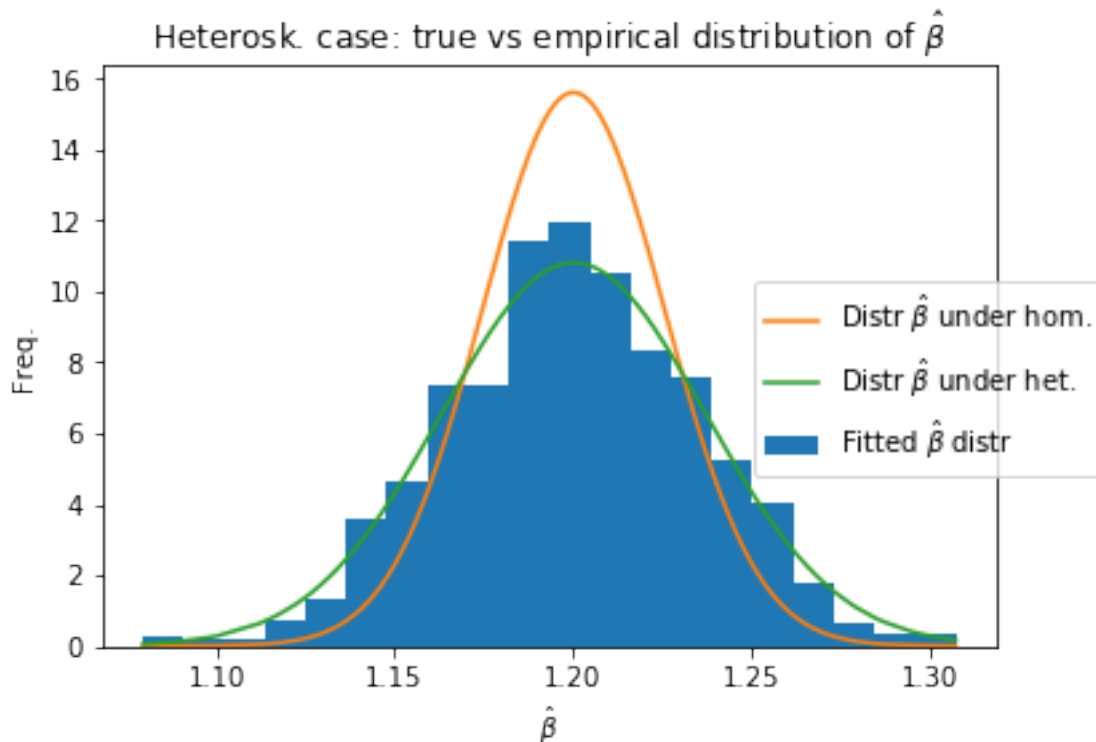
[13]: # 5) plot empirical vs true density of betahat
fig, ax = plt.subplots()
ax.hist(betas2, bins=20, density=True, label="Fitted  $\hat{\beta}$  distr" )
ax.plot(sorted(betas2), stats.norm.pdf(sorted(betas2), beta, betahat_se_hom2),
        label="Distr  $\hat{\beta}$  under hom.")
ax.plot(sorted(betas2), stats.norm.pdf(sorted(betas2), beta, betahat_se_het2),
        label="Distr  $\hat{\beta}$  under het.")
ax.set_xlabel(" $\hat{\beta}$ ")
ax.set_ylabel("Freq.")
ax.set_title("Heterosk. case: true vs empirical distribution of  $\hat{\beta}$ ")
fig.legend(loc="center right")

```

```

[13]: <matplotlib.legend.Legend at 0x2f3634c8668>

```



2.3 Serial correlation

2.3.1 Simulate data

We have a sample $(x_t, y_t)_{t=1}^T$ generated from the following model

$$y_t = \alpha + \beta x_t + e_t$$

where $x \perp e$, x follows a $U(0,1)$ and $e_t \sim MA(1)$ process. A $MA(1)$ process is simple a time series process of the form

$$e_t = u_t + \theta u_{t-1}$$

where u_t, u_{t-1} are white noise (i.e. orthogonal, zero-mean random variables).

```
[14]: from statsmodels.tsa.arima_process import arma_generate_sample # generate MA

# simulation parameters
npop = 1000
alpha = .5
beta = 1.2
sigma = .7 # MA(1) error sd
theta = .8 # MA(1) param
```

```

# generate random data
np.random.seed(1234) # important! Seed for reproducibility
x3 = np.random.uniform(0, 1, npop)
e3 = arma_generate_sample([1], [1, theta], npop, sigma)
y3 = alpha + beta*x3 + e3

# create dataframe with population data
df3 = pd.DataFrame()
df3['y'] = y3
df3['x'] = x3

```

```


[15]: fig, ax = plt.subplots()
      ax.set_title("Scatter of $y$ on $x$")
      ax.scatter(df3['x'], e3, s=4)
      ax.set_xlabel('$x$')
      ax.set_ylabel('$y$')

```

```

[15]: Text(0, 0.5, '$y$')

```



```

[16]: # 1) fit beta hat from full dataset under homosked
      reg_mod3 = smf.ols("y ~ x", df3)
      fit_hom3 = reg_mod3.fit() # homosked errors
      betahat_se_hom3 = fit_hom3.bse[1]

```

2.3.2 Check for serial correlation

```
[17]: from scipy.stats import norm # get quantile funct for N(mu, sigma)

# get resid from heterosk model
ehat = fit_hom3.resid # it does not matter if the original fit was assuming_
    →hom

# create ehat_t*ehat_{t-j} for j=0, ..., 4
def cross_prod_lag(x, k):
    x_std = (x - np.mean(x))/np.std(x)
    return x_std*x_std.shift(-k)
e2_lag0 = cross_prod_lag(ehat, 0)
e2_lag1 = cross_prod_lag(ehat, 1)
e2_lag2 = cross_prod_lag(ehat, 2)
e2_lag3 = cross_prod_lag(ehat, 3)
e2_lag4 = cross_prod_lag(ehat, 4)
e2_lag5 = cross_prod_lag(ehat, 5)

# put cross-products in dataset
df_e2_lags = pd.DataFrame()
df_e2_lags['e2_lag1'] = e2_lag1
df_e2_lags['e2_lag2'] = e2_lag2
df_e2_lags['e2_lag3'] = e2_lag3
df_e2_lags['e2_lag4'] = e2_lag4
df_e2_lags['e2_lag4'] = e2_lag4
df_e2_lags['e2_lag5'] = e2_lag5

# run regressions on a constant
eac_reg1 = smf.ols('e2_lag1 ~ 1', df_e2_lags).fit()
eac_reg2 = smf.ols('e2_lag2 ~ 1', df_e2_lags).fit()
eac_reg3 = smf.ols('e2_lag3 ~ 1', df_e2_lags).fit()
eac_reg4 = smf.ols('e2_lag4 ~ 1', df_e2_lags).fit()
eac_reg5 = smf.ols('e2_lag5 ~ 1', df_e2_lags).fit()

# get point est
corr_lag1 = eac_reg1.params[0]
corr_lag2 = eac_reg2.params[0]
corr_lag3 = eac_reg3.params[0]
corr_lag4 = eac_reg4.params[0]
corr_lag5 = eac_reg5.params[0]
corr_lags = (corr_lag1, corr_lag2,
             corr_lag3, corr_lag4, corr_lag5)

# get SE
se_corr_lag1 = eac_reg1.bse[0]
```

```

se_corr_lag2 = eac_reg2.bse[0]
se_corr_lag3 = eac_reg3.bse[0]
se_corr_lag4 = eac_reg4.bse[0]
se_corr_lag5 = eac_reg5.bse[0]
se_corr_lags = (se_corr_lag1, se_corr_lag2,
                se_corr_lag3, se_corr_lag4, se_corr_lag5)

# put corr, SE and lags in dataframe
df_acf_res = pd.DataFrame()
df_acf_res['lag'] = np.arange(6)[1:]
df_acf_res['corr'] = corr_lags
df_acf_res['se'] = se_corr_lags
df_acf_res['ci'] = df_acf_res['se']*norm.ppf(.995)
df_acf_res

```

```

[17]:
   lag    corr      se      ci
0    1  0.460620  0.035812  0.092246
1    2 -0.066885  0.031966  0.082338
2    3 -0.035929  0.031871  0.082093
3    4  0.049611  0.031822  0.081968
4    5  0.068970  0.030310  0.078074

```

```

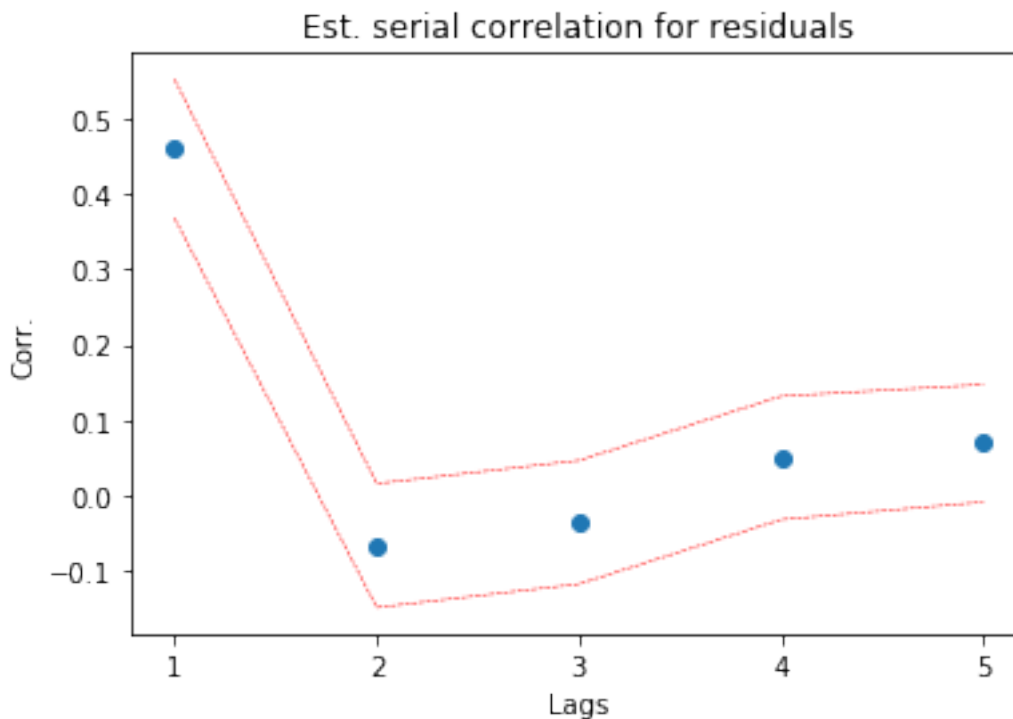
[18]: fig, ax = plt.subplots()
ax.scatter(df_acf_res['lag'], df_acf_res['corr'])
ax.plot(df_acf_res['lag'], df_acf_res['corr'] + df_acf_res['ci'],
        linestyle = '--', color = 'red', linewidth = .5)
ax.plot(df_acf_res['lag'], df_acf_res['corr'] - df_acf_res['ci'],
        linestyle = '--', color = 'red', linewidth = .5)
ax.set_xticks(df_acf_res['lag'])
ax.set_xlabel('Lags')
ax.set_ylabel('Corr.')
ax.set_title('Est. serial correlation for residuals')

```

```

[18]: Text(0.5, 1.0, 'Est. serial correlation for residuals')

```



2.3.3 True vs empirical distribution OLS estimator

Here is the experiment that we are going to run in a somewhat heuristic way:

1. we fit a linear regression $y \sim x$ on the whole sample and save the estimated $\hat{\beta}$ and its SE. We also (wrongly) assume *homoskedastic* errors.
2. we fit a linear regression $y \sim x$ on the whole sample and save the estimated $\hat{\beta}$ and its SE. We also (correctly) assume *serially correlated* errors.
3. we sample a random subset of our data $(x_t^{(i)}, y_t^{(i)})_t$
4. we fit the same linear regression $y \sim x$ on this sub-sample and save the estimated $\hat{\beta}^{(i)}$
5. we repeat steps 3-4) for $i = 1, 2, \dots, N_{sim}$ and obtain an empirical distribution of all the $(\hat{\beta}^{(i)})_{i=1}^N$

Then, if we correctly specified our model, we should expect that the empirical distribution of the $\hat{\beta}^{(i)}$ approaches the “true” asymptotic distribution of $\hat{\beta}$ which - under $MA(1)$ errors - we know is

$$\sqrt{n} \begin{pmatrix} \hat{\alpha} - \alpha \\ \hat{\beta} - \beta \end{pmatrix} \rightarrow \mathcal{N} \left(\begin{pmatrix} 0 \\ 0 \end{pmatrix}, (X'X)^{-1} X' \Omega X (X'X)^{-1} \right)$$

with

$$\Omega = \begin{pmatrix} (1+\theta^2)\sigma^2 & \theta\sigma^2 & 0 & \dots & 0 & 0 \\ \theta\sigma^2 & (1+\theta^2)\sigma^2 & \theta\sigma^2 & \dots & 0 & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & \dots & \theta\sigma^2 & (1+\theta^2)\sigma^2 \end{pmatrix}$$

```
[19]: # 2) fit beta hat from full dataset under serial correlation
nw_cov_opt = {'maxlags':1, 'use_correction':True}
fit_nw3 = reg_mod3.fit(cov_type = 'HAC', cov_kwds=nw_cov_opt) # NW SE
betahat_se_nw3 = fit_nw3.bse[1]
print(f"beta SE hom {betahat_se_hom3}")
print(f"beta SE NW {betahat_se_nw3}")
fit_nw3.summary()
```

beta SE hom 0.09470910054205718

beta SE NW 0.09912844755947796

```
[19]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                OLS Regression Results
=====
Dep. Variable:                  y      R-squared:                0.143
Model:                            OLS      Adj. R-squared:         0.142
Method:                 Least Squares      F-statistic:            151.6
Date:                Fri, 04 Oct 2019      Prob (F-statistic):      1.53e-32
Time:                  10:01:10      Log-Likelihood:         -1293.0
No. Observations:          1000      AIC:                   2590.
Df Residuals:              998      BIC:                   2600.
Df Model:                    1
Covariance Type:            HAC
=====

```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	0.5077	0.061	8.301	0.000	0.388	0.628
x	1.2207	0.099	12.315	0.000	1.026	1.415

```

=====
Omnibus:                 2.827      Durbin-Watson:           1.079
Prob(Omnibus):            0.243      Jarque-Bera (JB):        2.948
Skew:                    -0.049      Prob(JB):                0.229
Kurtosis:                 3.247      Cond. No.:               4.34
=====

```

Warnings:

```
[1] Standard Errors are heteroscedasticity and autocorrelation robust (HAC)
using 1 lags and with small sample correction
"""
```



```

[20]: # 2-3) fit model on random sub-sample
def bootstrap_lm_nw(x, params):
    """ Fit univariate lin reg to random sample

    The function assumes we know the true DGP

    Arguments
    -----
    x : np.array, univariate explanatory variable
    params : 1D np.array, alpha, beta, theta and sigma

    Returns
    -----
    beta_boot : estimated beta coeff on bootstrap sample
    """
    # generate bootstrap sample
    alpha, beta, theta, sigma = params
    n = len(x)
    indices = np.random.choice(np.arange(n), n)

    # generate residuals using right DGP
    x_boot = x[indices]
    e_boot = arma_generate_sample([1], [1, theta], n, sigma)
    y_boot = alpha + beta*x_boot + e_boot

    # create dataframe with bootstrapped data
    df = pd.DataFrame()
    df['y'] = y_boot
    df['x'] = x_boot

    # fit linear regression
    fit = smf.ols("y ~ x", df).fit()
    beta_boot = fit.params[1]

    return beta_boot

[21]: # fit on random sample
nsim = 1200
np.random.seed(3)
f3 = lambda i: bootstrap_lm_nw(df3['x'], [alpha, beta, theta, sigma])
betas3 = list(map(f3, range(nsim)))

[22]: # 5) plot empirical vs true density of betahat
fig, ax = plt.subplots()
ax.hist(betas3, bins=20, density=True, label="Fitted  $\hat{\beta}$  distr" )
ax.plot(sorted(betas3), stats.norm.pdf(sorted(betas3), beta, betahat_se_hom3),
        label="Distr  $\hat{\beta}$  under hom.")
ax.plot(sorted(betas3), stats.norm.pdf(sorted(betas3), beta, betahat_se_nw3),


```

```

        label="Distr  $\hat{\beta}$  under ser. corr.")
ax.set_xlabel(" $\hat{\beta}$ ")
ax.set_ylabel("Freq.")
ax.set_title("Serial corr. case: true vs empirical distribution of  $\hat{\beta}$ ")
fig.legend(loc="center right")

```

[22]: <matplotlib.legend.Legend at 0x2f363887668>



[]: